



COL380: Introduction to Parallel and Distributed Programming

Assignment-0 Report

Prof. Subodh Kumar

ARYAN SANTOSH SAHOO

2019CS10335

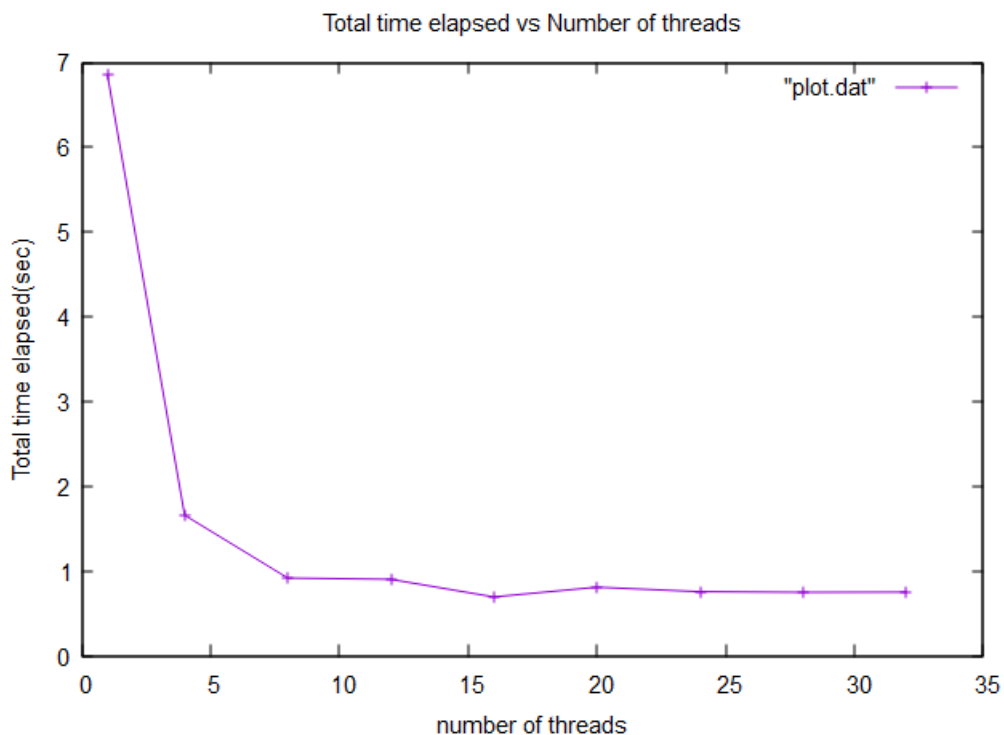
2.1 Perf Stat

Number of lines read: 1009072

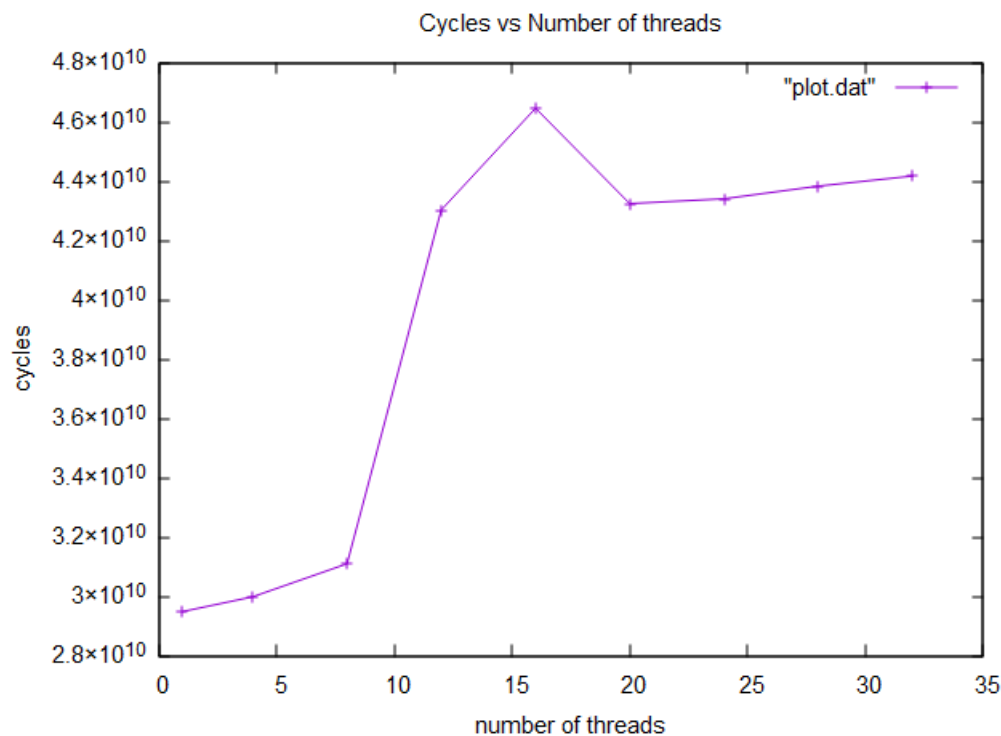
Number of runs: 6

Number of threads	Time elapsed (sec)	Cycles
1	6.853918571	29505554053
4	1.664863805	30006091171
8	0.923874249	31119619758
12	0.909170374	43036352278
16	0.701718816	46460936690
20	0.815588766	43250471275
24	0.764130756	43409226990
28	0.756588426	43839599972
32	0.758183896	44175903686

Time elapsed vs number of threads:



Cycles vs number of threads:



Perf stats for threads = 1

```

Performance counter stats for './classify rfile dfile 1009072 1':

      6,311.13 msec task-clock                    #    0.921 CPUs utilized
         158      context-switches              #    25.035 /sec
           8      cpu-migrations                 #     1.268 /sec
        15,347    page-faults                   #     2.432 K/sec
 29,50,55,54,053 cycles                        #     4.675 GHz
 64,86,21,58,042 instructions                  #     2.20  insn per cycle
 23,00,25,45,350 branches                      #     3.645 G/sec
   61,36,05,265  branch-misses                 #     2.67% of all branches
1,47,52,46,14,665 slots                       #    23.375 G/sec
 40,87,82,89,045 topdown-retiring              #    26.4% retiring
 65,95,21,80,673 topdown-bad-spec              #    42.7% bad speculation
 43,84,91,98,757 topdown-fe-bound              #    28.4% frontend bound
  3,91,09,64,866 topdown-be-bound              #     2.5% backend bound

 6.853918571 seconds time elapsed

 6.304360000 seconds user
 0.007975000 seconds sys

cs1190335@css7:~/Desktop/A0$

```

Analysis:

The above screenshot shows the perf stat command run on one thread to give the estimate of key CPU parameters like cycles, total time elapsed, branches and branch-misses.

As expected we see sharp decrease in the time elapsed as the number of threads increases but then it becomes almost constant as the number of threads increase above 8.

The reason for that being that each thread gets allotted different tasks and they together complete it in less time. But using too many threads can hurt performance as it adds overhead in scheduling. The code here in this case is not optimised and isn't very CPU extensive hence we see a usual decrease at first in the time elapsed and then a stagnation.

Branches misses also follows the same trend as increase in the number of threads means less misses as each thread has to take care of a smaller task now.

The number of cycles increases with number of threads because the context switching increases.

2.2 Perf Record

Perf record was run and a “perf.data” file was generated.

Perf report was run to inspect the file and then it was annotated for further analysis.

This “perf.data” was renamed to “perf_1.data”.

By seeing the annotated assembly code, we can point out the part of the code which is the most CPU time extensive.

38.03		jg	93
0.35		shl	\$0x6,%rax
0.06		add	%rbp,%rax
0.06	4e:	mov	%r13d,0x4(%r12)
0.38		mov	(%rax),%rdx
0.19		cmp	%r9d,0x8(%rax)
		↓ jbe	b9
		lea	(%rdx,%rdi,1),%rax
		add	%ebx,%ecx
2.50		mov	(%rax),%edx
0.04		add	\$0x1,%edx
0.08		mov	%edx,(%rax)
0.02		mov	%ecx,%eax
		cmp	%ecx,(%r8)
		↓ jbe	b0
0.00	70:	cltq	
0.02		lea	(%r10,%rax,8),%r12
		mov	0x8(%rsi),%eax
0.03		mov	(%r12),%edx
0.00		test	%eax,%eax
		↓ jle	a8
0.01		mov	(%rsi),%r11
		lea	-0x1(%rax),%r14d
		xor	%eax,%eax
0.00	8a:	mov	%eax,%r13d
9.26		cmp	(%r11,%rax,8),%edx
14.83		↑ jge	40
17.94	93:	→ lea	0x1(%rax),%r13
1.02		cmp	%rax,%r14
		↓ je	a8
		mov	%r13,%rax
14.86		↑ jmp	8a

Assembly instruction:

jg 93

(38.03% CPU time)

That line in the assembly code maps to the bool in “classify.h”

To get further clarity we change the “makefile” so that now it shows both the source code and assembly code, we do this by adding a “-g” flag.

This “perf.data” file was renamed to “perf_2_1.data”

```
37.53 | jg 93
0.36 | shl $0x6,%rax
0.06 | add %rbp,%rax
0.07 | 4e: _Z8classifyR4DataRK6Rangesj._omp_fn.0():
    | mov %r13d,0x4(%r12)
    | // and store the interval id in value. D is changed.
    | counts[v].increase(tid); // Found one key in interval v
0.35 | mov (%rax),%rdx
    | _ZN7Counter8increaseEj():
    | assert(id < _numcount);
0.17 | cmp %r9d,0x8(%rax)
    | ↓ jbe b9
    | _counts[id]++;
0.00 | lea (%rdx,%rdi,1),%rax
    | _Z8classifyR4DataRK6Rangesj._omp_fn.0():
    | for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop through all of Data
    | add %ebx,%ecx
    | _ZN7Counter8increaseEj():
2.34 | mov (%rax),%edx
0.04 | add $0x1,%edx
0.07 | mov %edx,(%rax)
    | _Z8classifyR4DataRK6Rangesj._omp_fn.0():
0.01 | mov %ecx,%eax
    | cmp %ecx,(%r8)
    | ↓ jbe b0
    | int v = D.data[i].value = R.range(D.data[i].key); // For each data, find the interval of data's key,
0.02 | 70: cltq
    | lea (%r10,%rax,8),%r12
    | _ZNK6Ranges5rangeEib():
    | if(strict) {
    | for(int r=0; r<_num; r++) // Look through all intervals
    | if(_ranges[r].strictlyin(val))
    | return r;
    | } else {
    | for(int r=0; r<_num; r++) // Look through all intervals
    | mov 0x8(%rsi),%eax
    | _Z8classifyR4DataRK6Rangesj._omp_fn.0():
0.02 | mov (%r12),%edx
    | _ZNK6Ranges5rangeEib():
    | test %eax,%eax
    | ↓ jle a8
    | if(_ranges[r].within(val))
0.02 | mov (%rsi),%r11
    | lea -0x1(%rax),%r14d
0.00 | xor %eax,%eax
    | 8a: mov %eax,%r13d
    | _ZNK5Range6withinEi():
    | return(lo <= val && val <= hi);
Press 'h' for help on key bindings
```

3. Hotspot Analysis

As stated above we had renamed the generated “perf.data” file to “perf_2_1.data” after changing the makefile to show the source code along with the assembly code.

Now we analyse the top hotspot in the code and see if it can be optimised to reduce runtime.

0.35	40:	cmp	0x4(%r11,%rax,8),%edx	Downloads	main.cpp	
37.53		jg	93	Music	Makefile	
0.36		shl	\$0x6,%rax	Pictures	rfile	
0.06		add	%rbp,%rax	Videos		
0.07	4e:	mov	%r13d,0x4(%r12)	Other Locations		
		// and store the interval id in value. D is changed.				
		counts[v].increase(tid); // Found one key in interval v				
0.35		mov	(%rax),%rdx			
		_ZN7Counter8increaseEj():				
		assert(id < _numcount);				
0.17		cmp	%r9d,0x8(%rax)			

0.02		mov	(%rsi),%r11	Downloads	main.cpp	
		lea	-0x1(%rax),%r14d	Music	Makefile	
0.00	8a:	xor	%eax,%eax	Pictures	rfile	
		mov	%eax,%r13d	Videos		
		_ZNK5Range6withinEi():				
		return(lo <= val && val <= hi);				
9.50		cmp	(%r11,%rax,8),%edx			
14.76		jge	40			
		_ZNK6Ranges5rangeEib():				
		for(int r=0; r<_num; r++) // Look through all intervals				
18.19	93:	lea	0x1(%rax),%r13			
0.99		cmp	%rax,%r14			
		je	a8			
0.00		mov	%r13,%rax			
15.11		jmp	8a			
		nop				
	a8:	mov	%rbp,%rax			
		return	r;			
		}				

We can see that the highlighted part of the assembly code takes the most CPU time.

Upon closer analysis we find out that it occurs because of a call from *"classify.cpp"* to *"classify.h"* where a certain variable is searched throughout the entire range using Boolean conditions, it is very inefficient and the performance can be improved by hashing or sorting through the range and searching by index.

For the final part to get an estimate of the branches, branch misses, cache misses, page faults and CPU cycles the following perf record command was run in the terminal:

```
perf record -e branch-instructions,branch-misses,cache-misses,page-faults,cpu-cycles make run
```

The generated *"perf.data"* file was then renamed to *"perf_2_2.data"*

4. Memory Profiling

Perf mem record was run and the data was stored into “perf_3.data”

This file contains data on the memory resources used by the code.

Two hotspots were identified:

1.20	<pre>_ZN7Counter8increaseEj(): assert(id < _numcount); cmp %r9d,0x8(%rax) ↓ jbe b9 _counts[id]++; lea (%rdx,%rdi,1),%rax _Z8classifyR4DataRK6Rangesj._omp_fn.0(): for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop through all of Data add %ebx,%ecx _ZN7Counter8increaseEj():</pre>	Downloads Music Pictures Videos Trash	main.cpp Makefile rfile
91.71	<pre> mov (%rax),%edx add \$0x1,%edx mov %edx,(%rax) _Z8classifyR4DataRK6Rangesj._omp_fn.0(): mov %ecx,%eax cmp %ecx,(%r8) ↓ jbe b0 int v = D.data[i].value = R.range(D.data[i].key); // For each data, find the interval of data's key, 70: cltq lea (%r10,%rax,8),%r12</pre>	+ Other Locations	
0.03			

0.09	<pre> xor %esi,%esi D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2. lea -0x4(%r14,%rdx,4),%r11 lea 0x8(%rcx),%rdx lea (%rdx,%r12,1),%r8 ↓ jmp 64 nop add \$0x8,%rdx</pre>	Downloads Music Pictures Videos Trash	main.cpp Makefile rfile
0.07	<pre>64: cmp %eax,0x4(%rcx) jne 81</pre>		
96.85	<pre>D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.</pre>		
0.59	<pre> mov 0x18(%rbx),%r9 mov (%rcx),%rcx mov %esi,%r10d add \$0x1,%esi</pre>		
0.49	<pre> add (%r11),%r10d mov 0x8(%r9),%r9 mov %rcx,(%r9,%r10,8)</pre>		
1.60	<pre>for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and 81: mov %rdx,%rcx cmp %rdx,%r8 jne 60</pre>		

Perf record -e cache-misses command was used to add the cache misses to the report as well this was stored in as “perf_5_1.data”

The code was modified to make it more cache friendly keeping the algorithm largely same.

Reducing the number of loops and making iterations a bit easier the improve in time complexity was approximately 10ms but the code has better cache handling.

Perf mem record was run again and data stored into “perf_4.data”

The hotspots were identified and they are now different from those before optimisation.

0.14	mov 0x8(%r13),%rdx	Downloads	main.cpp
	int rcount = 0;	Music	
	xor %esi,%esi	Pictures	Makefile
	lea 0x8(%rdx),%rax	Videos	rfile
	lea (%rax,%r12,1),%r8	Trash	
	↓ jmp cc		
	nop		
0.04	c8: add \$0x8,%rax		
	if(D.data[d].value == r){ // If the data item is in this interval		
0.70	cc: cmp %ecx,0x4(%rdx)		
97.85	jne eb		
	D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.		
0.32	mov 0x18(%rdi),%r10		
	mov (%rdx),%rdx		
	mov %esi,%r11d		
	add \$0x1,%esi		
0.50	add -0x4(%r9,%rcx,4),%r11d		
	mov 0x8(%r10),%r10		
	mov %rdx,(%r10,%r11,8)		
	for(int d=0; d<D.ndata; d++){ // For each interval, thread loops through all of data and		
0.37	eb: →mov %rax,%rdx		
	cmp %rax,%r8		
	↑ jne c8		

	bool within(int val) const { // Return if val is within this range	Starred	classify.cpp
0.02	return(lo <= val && val <= hi);	Home	classify.h
0.13	68: cmp 0x4(%rdi,%rdx,8),%esi	Documents	dnfile
	↓ jg b3	main.cpp	
	shl \$0x6,%rdx	Music	Makefile
	add %rbx,%rdx	Pictures	rfile
	_Z8classifyR4DataRK6Rangesj._omp_fn.0():	Videos	
	75: mov %r11d,0x4(%rcx)	Trash	
	// and store the interval id in value. D is changed.		
	counts[v].increase(tid); // Found one key in interval v		
0.54	mov (%rdx),%rsi		
	_ZN7Counter8increaseEj():		
	assert(id < _numcount);		
0.68	cmp 0x8(%rdx),%eax		
	↓ jae 168		
	_counts[id]++;		
	lea (%rsi,%r9,1),%rdx		
	add \$0x8,%rcx		
97.80	mov (%rdx),%esi		
	add \$0x1,%esi		
	mov %esi,(%rdx)		
	_Z8classifyR4DataRK6Rangesj._omp_fn.0():		
	for(int i=tid*len; i < (tid + 1)*len; i++) { // Threads together share-loop through all of Data		
	cmp %r10,%rcx		
	↓ je d0		
	_ZNK6Ranges5rangeEib():		
	if(strict) {		
	for(int r=0; r<num; r++) // Look through all intervals		
	if(_ranges[r].strictlyin(val))		
	return r;		
	} else {		

Perf record -e cache-misses command was used again and the data stored in *"perf_5_2.data"*